

Laboratory Exercises, C++ Programming

General information:

- The course has four compulsory laboratory exercises.
- You shall work in groups of two people. Sign up for the labs at `sam.cs.lth.se/Labs`.
- The labs are mostly homework. Before each lab session, you must have done the assignments (A1, A2, ...) in the lab, written and tested the programs, and so on. Reasonable attempts at solutions count; the lab assistant is the judge of what's reasonable. Contact a teacher if you have problems solving the assignments.
- Smaller problems with the assignments, e.g., details that do not function correctly, can be solved with the help of the lab assistant during the lab session.
- Extra labs are organized only for students who cannot attend a lab because of illness. Notify Roger Henriksson (`Roger.Henriksson@cs.lth.se`) if you fall ill, *before* the lab.

The labs are about:

1. Basic C++ programming, compiling, linking.
2. Introduction to the standard library.
3. Strings and streams.
4. Standard containers and algorithms.

Practical information:

- You will use many half-written "program skeletons" during the lab. You must download the necessary files from the course homepage before you start working on the lab assignments.
- The lab files are in separate directories `lab[1-4]` and are available in gzipped tar format. Download the tar file and unpack it like this:

```
tar xzf lab1.tar.gz
```

This creates a directory `lab1` in the current directory.

Good sources of information about C++:

- <http://www.cplusplus.com>
- <http://www.cppreference.com>

1 Basic C++ Programming, Compiling, Linking

Objective: to introduce C++ programming in a Unix environment.

Read:

- Book: basic C++, variables and types including pointers, expressions, statements, functions, simple classes, `ifstream`, `ofstream`.
- Manpages for `gcc`, `g++`, and `ld`.
- GNU make: <http://www.gnu.org/software/make/manual/>
- Valgrind: <http://www.valgrind.org>

1 Introduction

Different C++ compilers are available in a Unix environment, for example `g++` from GNU (see <http://gcc.gnu.org/>) and `clang++` from the Clang project (see <http://clang.llvm.org/>). The GNU Compiler Collection, GCC, includes compilers for many languages, the Clang collection only for “C-style” languages. `g++` and `clang++` are mostly compatible and used in the same way (same compiler options, etc.). In the remainder of the lab we mention only `g++`, but everything holds for `clang++` as well.

Actually, `g++` is not a compiler but a “driver” that invokes other programs:

Preprocessor (`cpp`): takes a C++ source file and handles preprocessor directives (`#include` files, `#define` macros, conditional compilation with `#ifdef`).

Compiler: the actual compiler that translates the input file into assembly language.

Assembler (`as`): translates the assembly code into machine code, which is stored in object files.

Linker (`ld`): collects object files into an executable file.

A C++ source code file is recognized by its extension. We will use `.cc`, which is the extension recommended by GNU. Another commonly used extension is `.cpp`.

In C++ (and in C) declarations are collected in header files with the extension `.h`. To distinguish C++ headers from C headers other extensions are sometimes used, such as `.hpp` or `.hh`. We will use `.h`.

A C++ program normally consists of many classes that are defined in separate files. It must be possible to compile the files separately. The program source code should be organized like this (a main program that uses a class `List`):

- Define the list class in a file `list.h`:

```
#ifndef LIST_H // include guard
#define LIST_H
// include necessary headers here

class List {
public:
    List();
    int size() const;
    ...
private:
    ...
};
#endif
```

- Define the class member functions in a file *list.cc*:

```
#include "list.h"
// include other necessary headers

List::List() { ... }
int List::size() const { ... }
...
```

- Define the main function in a file *ltest.cc*:

```
#include "list.h"
#include <iostream>

int main() {
    List list;
    std::cout << "Size: " << list.size() << std::endl;
    ...
}
```

The include guard is necessary to prevent multiple definitions of names. Do *not* write function definitions in a header file (except inline functions and template functions).

The g++ command line looks like this:

```
g++ [options] [-o outfile] infile1 [infile2 ...]
```

The .cc files are compiled separately. The resulting object files (.o files) are linked into an executable file *ltest*, which is then executed:

```
g++ -c list.cc -std=c++11
g++ -c ltest.cc -std=c++11
g++ -o ltest ltest.o list.o
./ltest
```

The -c option directs the driver to stop before the linking phase and produce an object file, named as the source file but with the extension .o instead of .cc.

A1. Write a "Hello, world!" program in a file *hello.cc*, compile and test it.

2 Options and messages

There are more options to the g++ command than were mentioned in section 1. Your source files must compile correctly using the following command line:

```
g++ -c -O2 -Wall -Wextra -pedantic-errors -Wold-style-cast -std=c++11 file.cc
```

Short explanations (you can read more about these and other options on the gcc and g++ manpages):

<code>-c</code>	just produce object code, do not link
<code>-O2</code>	optimize the object code (perform nearly all supported optimizations)
<code>-Wall</code>	print most warnings
<code>-Wextra</code>	print extra warnings
<code>-pedantic-errors</code>	treat “serious” warnings as errors
<code>-Wold-style-cast</code>	warn for old-style casts, e.g., <code>(int)</code> instead of <code>static_cast<int></code>
<code>-std=c++11</code>	follow the new C++ standard (use <code>-std=c++0x</code> on early versions of g++)
<code>-stdlib=libc++</code>	Clang only — use Clang’s own standard library instead of GNU’s <code>libstdc++</code>

Do not disregard warning messages. Even though the compiler chooses to “only” issue warnings, your program is erroneous or at least questionable.

Some of the warning messages are produced by the optimizer and will therefore not be output if the `-O2` flag is not used. But you must be aware that optimization takes time, and on a slow machine you may wish to remove this flag during development to save compilation time. Some platforms define higher optimization levels, `-O3`, `-O4`, etc. You should not use these optimization levels unless you know very well what their implications are.

It is important that you become used to reading and understanding the GCC error messages. The messages are sometimes long and may be difficult to understand, especially when the errors involve the standard library template classes (or any other complex template classes).

3 Introduction to make

You have to type a lot in order to compile and link C++ programs — the command lines are long, and it is easy to forget an option or two. You also have to remember to recompile all files that depend on a file that you have modified.

There are tools that make it easier to compile and link, “build”, programs. These may be integrated development environments (Eclipse, Visual Studio, ...) or separate command line tools. In Unix, *make* is the most important tool. Make works like this:

- it reads a “makefile” when it is invoked. Usually, the makefile is named *Makefile*.
- The makefile contains a description of dependencies between files (which files that must be recompiled/relinked if a file is updated).
- The makefile also contains a description of how to perform the compilation/linking.

As an example, we take the list program from section 1. The files *list.cc* and *ltest.cc* must be compiled and then linked. Instead of typing the command lines, you just enter the command *make*. Make reads the makefile and executes the necessary commands.

A minimal makefile, without all the compiler options, looks like this:

```
# The following rule means: "if ltest is older than ltest.o or list.o,
# then link ltest".
ltest: ltest.o list.o
    g++ -o ltest ltest.o list.o

# Rules to create the object files.
ltest.o: ltest.cc list.h
    g++ -c ltest.cc -std=c++11
list.o: list.cc list.h
    g++ -c list.cc -std=c++11
```

A rule specifies how a file (the *target*), which is to be generated, depends on other files (the *prerequisites*). The line following the rule contains a shell command, a *recipe*, that generates the target. The recipe is executed if any of the prerequisites are older than the target. It must be preceded by a TAB character, *not* eight spaces.

A2. The file *Makefile* in the *lab1* directory contains the makefile described above. The files *list.h*, *list.cc*, and *ltest.cc* are in the same directory. Experiment:

Run `make`. Run `make` again. Delete the executable program and run `make` again. Change one or more of the source files (it is sufficient to touch them) and see what happens. Run `make ltest.o`. Run `make notarget`. Read the manpage and try other options.

4 More Advanced Makefiles

4.1 Implicit Rules

Make has *implicit rules* for many common tasks, for example producing *.o*-files from *.cc*-files. The recipe for this task is:

```
$(CXX) $(CPPFLAGS) $(CXXFLAGS) -c -o $@ $<
```

CXX, *CPPFLAGS*, and *CXXFLAGS* are variables that the user can define. The expression `$(VARIABLE)` evaluates a variable, returning its value. *CXX* is the name of the C++ compiler, *CPPFLAGS* are the options to the preprocessor, *CXXFLAGS* are the options to the compiler. `$@` expands to the name of the target, `$<` expands to the first of the prerequisites.

There is also an implicit rule for linking, where the recipe (after some variable expansions) looks like this:

```
$(CC) $(LDFLAGS) $^ $(LOADLIBES) $(LDLIBS) -o $@
```

LDFLAGS are options to the linker, such as `-Ldirectory`. *LOADLIBES* and *LDLIBS*¹ are variables intended to contain libraries, such as `-llab1`. `$^` expands to all prerequisites. So this is a good rule, except for one thing: it uses `$(CC)` to link, and *CC* is by default the C compiler `gcc`, not `g++`. But if you change the definition of *CC*, the implicit rule works also for C++:

```
# Define the linker
CC = g++
```

4.2 Phony Targets

Make by default creates the first target that it finds in the makefile. By convention, the first target should be named *all*, and it should make all the targets. But suppose that a file *all* is created in the directory that contains the makefile. If that file is newer than the *ltest* file, a make invocation will do nothing but say `make: Nothing to be done for 'all'.`, which is not the desired behavior. The solution is to specify the target *all* as a *phony target*, like this:

```
all: ltest
.PHONY: all
```

¹ There doesn't seem to be any difference between *LOADLIBES* and *LDLIBS* — they always appear together and are concatenated. Use *LDLIBS*.

Another common phony target is *clean*. Its purpose is to remove intermediate files, such as object files, and it has no prerequisites. It typically looks like this:

```
.PHONY: clean
clean:
    rm -f *.o ltest
```

4.3 Generating Prerequisites Automatically

While you're working with a project the prerequisites are often changed. New `#include` directives are added and others are removed. In order for make to have correct information about the dependencies, the makefile must be modified accordingly. This is a tedious task, and it is easy to forget a dependency.

The C++ preprocessor can be used to generate prerequisites automatically. The option `-MMD`² makes the preprocessor look at all `#include` directives and produce a file with the extension `.d` which contains the corresponding prerequisite. Suppose the file *ltest.cc* contains the following `#include` directive:

```
#include "list.h"
```

The compiler produces a file *ltest.d* with the following contents:

```
ltest.o : ltest.cc list.h
```

The `.d` files are included in the makefile, so it functions the same way as if we had written the rules ourselves.

4.4 Putting It All Together

The makefile below can be used as a template for makefiles in many (small) projects. To add a new target you must:

1. add the name of the executable to the definition of `PROGS`,
2. add a rule which specifies the object files that are necessary to produce the executable.

See the make manual, section 4.14, if you are interested in details about the `%.d: %.cc` rule.

```
# Define the compiler and the linker. The linker must be defined since
# the implicit rule for linking uses CC as the linker. g++ can be
# changed to clang++.
CXX = g++
CC  = $(CXX)

# Generate dependencies in *.d files
DEPFLAGS = -MT $@ -MMD -MP -MF $*.d

# Define preprocessor, compiler, and linker flags. Uncomment the # lines
# if you use clang++ and wish to use libc++ instead of GNU's libstdc++.
# -g is for debugging.
CPPFLAGS = -std=c++11
CXXFLAGS = -O2 -Wall -Wextra -pedantic-errors -Wold-style-cast
CXXFLAGS += -std=c++11
CXXFLAGS += -g
CXXFLAGS += $(DEPFLAGS)
LDFLAGS = -g
```

² The option `-MMD` generates prerequisites as a side effect of compilation. If you only want the preprocessing but no compilation, `-MM` can be used.

```

#CPPFLAGS = -stdlib=libc++
#CXXFLAGS += -stdlib=libc++
#LDFLAGS += -stdlib=libc++

# Targets
PROGS = ltest

all: $(PROGS)

# Targets rely on implicit rules for compiling and linking
ltest: ltest.o list.o

# Phony targets
.PHONY: all clean

# Standard clean
clean:
    rm -f *.o $(PROGS)

# Include the *.d files
SRC = $(wildcard *.cc)
-include $(SRC:.cc=.d)

```

- A3. The better makefile is in the file *BetterMakefile*. Rename this file to *Makefile*, experiment. The first time you run make you will get warnings about *.d*-files that don't exist. This is normal. Also, the compiler will warn about unused parameters. These warnings will disappear when you implement the member functions.

Look at the generated *.d* files. Add a rule to build your "Hello world!" program.

5 A List Class

- A4. The class `List` describes a linked list of integers.³ The numbers are stored in nodes. A node has a pointer to the next node (`nullptr` in the last node).

In this assignment you shall use raw pointers and manual memory allocation and deletion. This is common in "library classes" which must be very efficient and are assumed to be error free. In an application you would use one of the safe pointer types that were introduced in the new standard.

```

class List {
public:
    /* creates an empty list */
    List();

    /* destroys this list */
    ~List();

    /* returns true if d is in the list */
    bool exists(int d) const;

    /* returns the size of the list */
    int size() const;

    /* returns true if the list is empty */
    bool empty() const;

    /* inserts d into this list as the first element */

```

³ In practice, you would never write your own list class. There are several alternatives in the standard library.


```

void insertFirst(int d);

/* removes the first element less than/equal to/greater than d,
   depending on the value of df. Does nothing if there is no value
   to remove. The enum values are accessed with List::DeleteFlag::LESS,
   ..., outside the class */
enum class DeleteFlag { LESS, EQUAL, GREATER };
void remove(int d, DeleteFlag df = DeleteFlag::EQUAL);

/* prints the contents of this list */
void print() const;

/* forbid copying of lists */
List(const List&) = delete;
List& operator=(const List&) = delete;
private:
/* a list node */
struct Node {
    int value; // the node value
    Node* next; // pointer to the next node, nullptr in the last node
    Node(int v, Node* n) : value(v), next(n) {}
};

Node* first; // pointer to the first node
};

```

Node is a struct, i.e., a class where the members are public by default. This is not dangerous, since Node is private to the class.

The copy constructor and assignment operator are deleted, so lists cannot be copied.

Implement the member functions in *list.cc*, build and test. Execution errors like “segmentation fault” are addressing errors. Read section 6 about finding execution errors.

A5. Implement a class Coding with two static methods:

```

/* For any character c, encode(c) is a character different from c */
static unsigned char encode(unsigned char c);

/* For any character c, decode(encode(c)) == c */
static unsigned char decode(unsigned char c);

```

Use a simple method for coding and decoding. Then write a program, *encode*, that reads a text file⁴, encodes it, and writes the encoded text to another file. The command line:

```
./encode file
```

should run the program, *encode file*, and write the output to *file.enc*.

Write another program, *decode*, that reads an encoded file, decodes it, and writes the decoded text to another file *file.dec*. The command line should be similar to that of the *encode* program. Add rules to the makefile for building the programs.

Test your programs and check that a file that is first encoded and then decoded is identical to the original. Use the Unix *diff* command.

Note: the programs will work also for files that are UTF-8 encoded. In UTF-8 characters outside the “ASCII range” are encoded in two bytes, and the *encode* and *decode* functions will be called twice for each such character.

⁴ Note that you cannot use `while (infile >> ch)` to read all characters in *infile*, since `>>` skips whitespace. Use `infile.get(ch)` instead. Output with `outfile << ch` should be ok, but `outfile.put(ch)` looks more symmetric.

6 Finding Errors

6.1 Debugging

With the GNU debugger, *gdb*, you can control a running program (step through the program, set breakpoints, inspect variable values, etc.). Debug information is inserted into the executable program when you compile and link with the option `-g`. Preferably you should also turn off optimization (no `-O2` option). From *g++* version 4.8 there is a new option `-Og`, which turns on all optimizations that do not conflict with debugging.

A program is executed under control of *gdb* like this:

```
gdb ./ltest
```

Some useful commands:

<code>help [command]</code>	Get help about <i>gdb</i> commands.
<code>run [args...]</code>	Run the program (with arguments).
<code>continue</code>	Continue execution.
<code>next</code>	Step to the next line <i>over</i> function calls.
<code>step</code>	Step to the next line <i>into</i> function calls.
<code>where</code>	Print the call stack.
<code>list [nbr]</code>	List 10 lines around the current line or around line <i>nbr</i> (the following lines if repeated).
<code>break func</code>	Set a breakpoint on the first line of a function <i>func</i> .
<code>break nbr</code>	Set a breakpoint at line <i>nbr</i> in the current file.
<code>print expr</code>	Print the value of the expression <i>expr</i> .
<code>watch var</code>	Set a watchpoint, i.e., watch all modifications of a variable. Can be very slow but can be the best solution to find some bugs.

A6. (Optional) Run *ltest* under control of *gdb*, try the commands.

6.2 Memory-Related Errors

In Java, many errors are caught by the compiler (use of uninitialized variables) or by the runtime system (addressing outside array bounds, dereferencing null pointers, etc.). In C++, errors of this kind are not caught, instead they result in erroneous results or faults during program execution. Furthermore, you get no information about where in the program the error occurred. Since deallocation of dynamic memory in C++ is manual, you also have a whole new class of errors (dangling pointers, double delete, memory leaks).

Valgrind is a tool (available under Linux and Mac OS X) that helps you find memory-related errors at the precise locations at which they occur. It does this by emulating an x86 processor and supplying each data bit with information about the usage of the bit. This results in slower program execution, but this is more than compensated for by the reduced time spent in searching for bugs.

Valgrind is easy to use. Compile and link as usual, then execute like this:

```
valgrind ./ltest
```

When an error occurs, you get an error message and a stack trace (and a lot of other information). At the end of execution, *valgrind* prints a “leak summary” which indicates the amount of dynamic memory that hasn’t been freed.

A7. Run *ltest* under control of *valgrind*. The leak summary should show that 0 bytes have been lost. If it doesn’t, the *List* destructor probably contains an error.

Introduce an addressing error in one of the `List` member functions (e.g., remove the check for end-of-list in `exists`). Run the program, first as usual, then under `valgrind`.

Introduce an error in the `List` destructor (e.g., delete all nodes but one). Run the program, first as usual, then under `valgrind`. Remove the errors that you introduced before continuing.

7 Object Code Libraries

A lot of software is shipped in the form of libraries, e.g., class packages. In order to use a library, a developer does not need the source code, only the object files and the headers. Object file libraries may contain thousands of files and cannot reasonably be shipped as separate files. Instead, the files are collected into library files that are directly usable by the linker.

7.1 Static Libraries

The simplest kind of library is a *static library*. The linker treats the object files in a static library in the same way as other object files, i.e., all code is linked into the executable files. In Unix, a static library is an archive file, `lib<name>.a`. In addition to the object files, an archive contains an index of the symbols that are defined in the object files.

A collection of object files `f1.o`, `f2.o`, `f3.o`, ..., are collected into a library `libfoo.a` using the `ar` command:

```
ar crv libfoo.a f1.o f2.o f3.o ...
```

(Some Unix versions require that you also create the symbol table with `ranlib libfoo.a`.) In order to link a program `main.o` with the object files `obj1.o`, `obj2.o` and with object files from the library `libfoo.a`, you use the following command line:

```
g++ -o main main.o obj1.o obj2.o -L. -lfoo
```

The linker searches for libraries in certain system directories. The `-L.` option makes the linker search also in the current directory.⁵ The library name (without `lib` and `.a`) is given after `-l`.

A8. Collect the object files `list.o` and `coding.o` in a library `liblab1.a`. Change the makefile so the programs (`ltest`, `encode`, `decode`) are linked with the library. The `-L` option belongs in `LD_FLAGS`, the `-l` option in `LD_LIBS`.

7.2 Shared Libraries

Since most programs use large amounts of code from libraries, executable files can grow very large. Instead of linking library code into each executable that needs it the code can be loaded at runtime. The object files should then be in *shared libraries*. When linking programs with shared libraries, the files from the library are not actually linked into the executable. Instead a “pointer” is established from the program to the library.

In Unix shared library files are named `lib<name>.so[.x.y.z]` (`.so` for shared objects, `.x.y.z` is an optional version number). The linker uses the environment variable `LD_LIBRARY_PATH` as the search path for shared libraries. In Microsoft Windows shared libraries are known as DLL files (dynamically loadable libraries).

⁵ You may have several `-L` and `-l` options on a command line. Example, where the current directory and the directory `/usr/local/mylib` are searched for the libraries `libfoo1.a` and `libfoo2.a`:

```
g++ -o main main.o obj1.o obj2.o -L. -L/usr/local/mylib -lfoo1 -lfoo2
```

- A9.** (Advanced, optional) Create a shared library with the object files *list.o* and *coding.o*. Link the executables using the shared library. Make sure they run correctly. Compare the sizes of the dynamically linked executables to the statically linked (there will not be a big difference, since the library files are small).

Use the command `ldd` (list dynamic dependencies) to inspect the linkage of your programs. Shared libraries are created by the linker, not the `ar` archiver. Use the `gcc` and `ld` manpages (and, if needed, other manpages) to explain the following sequence of operations:

```
g++ -fPIC -std=c++11 -c *.cc
g++ -shared -Wl,-soname,liblab1.so.1 -o liblab1.so.1.0 list.o coding.o
ln -s liblab1.so.1.0 liblab1.so.1
ln -s liblab1.so.1 liblab1.so
```

You then link with `-L. -llab1` as before. The linker merely checks that all referenced symbols are in the shared library. Before you execute the program, you must define `LD_LIBRARY_PATH` so it includes the current directory. You do this with the following command (on the command line):

```
export LD_LIBRARY_PATH=.:$LD_LIBRARY_PATH
```